

COMPUTATIONAL PHILOLOGY · CORPORA · NEGATIVE RESULTS

Reading Steppe Names in Chinese Transcription

Three corpora, two medieval lexicons, and a measured answer to a fifty-year-old question about how much chance a sound law may tolerate.

Aziz M. Uluğ, PhD

Cortechs.ai · aulug@cortechs.ai

Preprint. Data and code: github.com/ulug2004/chinese-transcription-channel

ABSTRACT

Chinese histories record hundreds of Inner Asian names (Xiongnu rulers, Turkic titles, tribal designations) in characters chosen for their sound. Two thousand years of change in Chinese pronunciation have left the modern Mandarin readings of those characters unrelated to what the scribes heard. This paper models the transcription process itself as a noisy channel, trains it where both sides are known, and measures what can and cannot be recovered.

We generated the following. **Three transcription corpora**, two of them extracted from sources not previously available as data: 9,329 Chinese–Mongolian pairs from the *Secret History of the Mongols*, 594 Chinese–Uyghur pairs from Ligeti's edition of a Ming glossary, and 1,017 verified Chinese–Sanskrit pairs separated out of a 27,110-entry Buddhist dictionary that mixes phonetic spellings with translations. **Two medieval Turkic lexicons** in machine-readable form, for which we found no downloadable dataset: 6,820 glossed headwords from the Türk Dil Kurumu index to Kāşgarî's *Dīwān Lughāt al-Turk* (1072–77), and 2,110 from Kuun's 1880 edition of the *Codex Cumanicus*. **A working channel model**, which on held-out data reconstructs 31% of Turkic words exactly and 84% within two segments. **Two positive readings** that survive every test we can apply. 若鞮, the element closing the formal names of six later chanyu, reads as *Inakt* "the faithful ones": the *Hou Hanshu* glosses it 孝 "filial", the *Codex Cumanicus* glosses *inak* as *fidelis*, and the plural in -t is an Old Turkic formation used with titles. 冒頓, the founder of the confederation, reads as *Bögü Tın* "sage spirit" on the tightest frame in the record, against a standard etymology that fails the one test the corpus decides cleanly. And **a set of measured negative results**: language attribution from transcription evidence is not supportable on available reference data, and full reconstruction of the Xiongnu names is not supportable at the density of evidence that survives: 1.4 observations per character, against the 60 that made the same method work on Mongolian.

The measurements also answer a question the field has carried unresolved since 1973. Doerfer argued that chance resemblance is pervasive in long-range comparison; his reviewer sharpened it into a threshold problem: *is it a sound law if 49% of cases show one correspondence and 51% another?* He proposed three classes that nobody could then apply, for want of numbers. Section 7 applies them, with confidence intervals, and finds that the threshold cannot be a constant: it is the baseline of the component being measured, and it moves with how many outcomes that component has.

Keywords: historical phonology · noisy-channel models · corpus construction · Xiongnu · Old Turkic · null models · negative results

1 The problem

When a Chinese scribe of the Han dynasty wrote down a foreign name, he chose characters for their sound. 冒頓, the founder of the Xiongnu state, is read *Mòdùn* in modern Mandarin, but Mandarin is not the language the characters were chosen in. In the Later Han period those two characters were pronounced approximately *mək tuən*.²² Every argument about what such a name meant, and about who the people bearing it were, depends on getting from the modern reading back to the period one.

That is a recoverable step: reconstructions of Later Han Chinese are published and are not in dispute here. The harder question is the one after it. Given a period reading, can the original foreign word be recovered? The transcription is lossy: Chinese had no *r*-, a restricted set of syllable shapes, and no way to write many foreign sounds, so the mapping from source word to spelling is many-to-one and its inverse is underdetermined. This is a noisy-channel problem, and it can be treated as one.

The interest of the Xiongnu case is that the answer is unknown and contested: proposals put the ruling house's language in the Turkic, Mongolic, Yeniseian and Iranian families, each resting on a handful of transcribed names. The interest of the *method* is that its reliability can be measured on cases where the answer is known, which is what most of this paper does.

2 A question the field has already asked

Nothing in the approach below is methodologically new, and it is worth saying so at the outset. Gerhard Doerfer made the argument in 1973, in *Lautgesetz und Zufall* ("Sound Law and Chance"), a study of what he called *omnicomparatism*: relating languages across great distances on the strength of resemblances.⁵ Chance produces such resemblances in quantity, was Doerfer's case, and comparativists have a repertoire of devices (he called them tricks) for making weak matches look strong. He included deliberately absurd comparisons of his own to show how easily the devices work.

A review in *Anthropos* the following year pressed the point into a question that can be answered with numbers.¹⁵ Sound laws, the reviewer observed, are not deterministic but probabilistic, and the literature never says how probabilistic they are permitted to be. *Is it a sound law if 49% of cases show the same correspondence and 51% show something else?* Where does regularity stop? He proposed three classes: deterministic laws, holding without exception; strong probabilistic laws, where one correspondence occurs in more than half of cases; and weak probabilistic laws, where one correspondence is merely the most common, and could apply none of them, for want of counts.

He also named the mathematics that was already available: the matching problem of Montmort (1708), which asks how often two shuffled sequences agree by accident. The permutation nulls used throughout this paper are that problem solved by simulation rather than by formula. What is new here is not the method but the measurement: we compute the numbers on this material instead of asserting them, which is what Doerfer was asking for.

3 Resources

A learned model needs examples, and for this problem almost none existed as data. We built five datasets: three of transcription pairs, and two lexicons of Turkic itself. We release or make reproducible all of them; §12 sets out which is which and why.

3.1 Corpus 1: *Secret History of the Mongols*

The *Secret History* is a thirteenth-century Mongolian chronicle preserved only in Chinese characters used phonetically, with an interlinear gloss. It is the largest known parallel text of its kind, and it had not been extracted as aligned pairs. Parsing the digital edition yields **9,329 Chinese–Mongolian transcription pairs**.

One detail decided whether the extraction worked. In the source encoding, a Chinese character followed by U+180B (a Mongolian free variation selector) marks an *annotation*, not part of the transcription. Treating those as data corrupts roughly a fifth of the corpus; the rule is one line, and finding it took considerably longer. Length agreement between the Chinese and Mongolian sides is 93.4%, which is the internal check that the alignment is real.

3.2 Corpus 2: Ligeti's Sino-Uyghur vocabulary

Ligeti published a Ming-dynasty Chinese–Uyghur glossary in two instalments in *Acta Orientalia*.^{17,18} The Turkic forms are given in Latin transcription and the Chinese in EFEO romanisation, which makes the text parseable in principle and awkward in practice. Extraction yields **594 Turkic–Chinese pairs**.

We had to handle four properties of the typography: the apostrophe is U+2019 rather than ASCII, Turkic forms may contain spaces, footnote digits intervene before the gloss, and the Turkic form is not always at the start of a line. Our first pass missed them and returned 297 pairs, half of them wrong. A further trap was structural rather than typographic: journal running heads carry the author's name on verso pages, so a filter written for the article title alone silently discarded half the article.

3.3 Corpus 3: verified Buddhist Sanskrit transcriptions

Buddhist translators into Chinese used two devices: 音譯, spelling a Sanskrit word by sound, and 意譯, translating its meaning. Only the first is transcription. A 27,110-entry terminology dictionary mixes them without marking which is which, and the two are not separable by inspection.

Two approaches failed before one worked. Character-set membership admits 福慧 for *punya*, which is a calque. Exclusivity, treating a character as phonetic only if it never appears in a translation, collapsed the corpus to 67 pairs. What worked was phonetic verification: score each candidate against the Later Han reading of its characters and keep only those whose sound actually matches. The result is **1,017 verified pairs**, and we keep the 7,480 rejected entries separately so anyone can audit the decision. The induced 音譯 inventory matches the standard akṣara table, which is an external check that the filter found something real.

3.4 Lexicons 1 and 2: Kāṣṣarī and the *Codex Cumanicus*

The tests in §6 need something the three corpora are not: dictionaries of Turkic itself, from near the period, against which a reading can be looked up. We found no downloadable dataset of either, so we extracted both.

Kāṣṣarī Mahmud's *Dīwān Lughāt al-Turk* (1072–77) is the earliest dictionary of any Turkic language.¹² The Türk Dil Kurumu index volume has a machine-readable text layer; parsing it yields **6,820 headwords** with Kāṣṣarī's Turkish glosses and his volume-and-page references.

The *Codex Cumanicus* (c. 1300) records Kipchak Turkic, compiled by Italian merchants and German missionaries, outsiders writing down a language they did not speak, which is the position the Chinese scribes were in. Géza Kuun's 1880 edition prints a Cumano-Latin vocabulary; parsing it yields **2,110 headwords** with Latin glosses.¹⁴ The Persian and German vocabularies that follow in the same volume we deliberately skipped: the Persian one would inject Iranian vocabulary into a Turkic lexicon and inflate every match rate computed from it.

4 Method

The channel is a probabilistic model of the scribe: for a source sound, what Chinese spelling would he choose? Character-to-sound alignments are learned by expectation–maximisation over the training pairs, without hand-alignment. Emission tables are *reading-anchored* with backoff: full reading, then onset plus rime, then onset, then a global distribution, so a character never seen in training still receives a distribution rather than nothing.

Decoding runs the channel backwards: given the characters, score candidate source words. The binding variable throughout is not corpus size but **observations per unit**: how many times the model has seen each thing it must estimate. That figure, not the number of pairs, is what predicts whether reconstruction works.

We pair every comparison in this paper with a null. For lexicon searches we build pseudo-names by resampling Later Han syllables from the same names being tested, so they carry the right shape and the right sound inventory and no meaning at all; the identical search is run on them. For language comparisons we subsample the corpora to equal size before scoring, because otherwise the thinner model wins by being vague. Negative controls, meaning the method run on cases whose answer is already known, are what exposed the failures reported in §6.

5 What the channel recovers

Every figure below is on held-out data. "Exact" means the model reconstructed the whole word correctly; "within 2" means it got no more than two segments wrong.

TRAINED ON → TESTED ON	EXACT	WITHIN 1	WITHIN 2
Mongolian → Mongolian (unseen chapters)	36.9%	66.3%	87.5%
Turkic → Turkic	30.6%	68.2%	83.5%
Sanskrit → Sanskrit	25.0%	57.9%	77.0%
Sanskrit → <i>Turkic</i> (wrong training language)	3.1%	6.9%	27.5%

TABLE 1 Each row is one experiment. Row 1 trains on some chapters of the *Secret History* and tests on chapters held back entirely, a document-level split, so no word from a test chapter influenced training. Rows 2 and 3 do the same within the Turkic and Sanskrit corpora, but those are glossaries without document structure, so the held-out portion is a random sample of pairs. Row 4 is the diagnostic: the channel is trained entirely on Sanskrit and asked to read Turkic. A model trained on Buddhist transcriptions barely functions on Turkic ones, 3.1% against 30.6%, which is the clearest evidence that the channel learns something specific to the language it was trained on rather than a generic Chinese-to-anything mapping.

CORPUS	PAIRS	OBSERVATIONS PER CHARACTER
Mongolian	9,329	~60
Sanskrit	1,017	~12
Turkic	594	~7
Xiongnu names	39 items	1.4

TABLE 2 The Xiongnu material is 112 character-tokens over 81 distinct characters. This is the number that governs everything in §6: the method does not fail on Xiongnu because Xiongnu is mysterious, but because 1.4 observations per character is not enough to estimate anything. The comparison is with 60 for the corpus where the same method reaches 36.9% exact.

6 What the evidence cannot support

Two claims commonly made from this kind of material do not hold up, and establishing that carefully is important.

6.1 Determining which language a name belongs to

The obvious application is settling whether the Xiongnu ruling house spoke Turkic, Yeniseian, Mongolic or something else. We built the comparison properly: score how well each candidate language explains the transcription, equalise corpus sizes, compare by Bayes factor.

It failed its own controls. Given Ming Uyghur transcriptions whose language is known with certainty, the method identified Turkic correctly in 69% of trials against a 50% chance baseline, and in 2 of 24 experimental cells it fell *below* chance. A diagnostic showed why: the phonotactic priors available for these families are close enough to each other that the comparison is largely measuring the reference data rather than the names. A procedure that cannot reliably identify a language it has been given cannot be trusted to identify one nobody has.

6.2 Reconstructing the Xiongnu names in full

At 1.4 observations per character the decoder returns candidates, but they are not constrained by evidence. An intermediate version of this work gated its output on in-domain accuracy and would have printed confident readings (撑犁 as *ara*, *kara*, *tara*) where true out-of-domain reliability was 3.1%. Re-gating on out-of-domain accuracy removed them. This was the closest the project came to publishing a wrong answer, and we record the correction here because the failure mode is general: a decoder always returns its best candidate, and best is not the same as supported.

6.3 Matching readings to dictionary words

There is an obvious next step, and it is the one every reader wants to take. Once 稽粥 is written out as *kei çuk*, it is hard not to hear Turkish *küçük* "small, younger". The method is old: look the reading up in a dictionary and see what it hits.

The trouble is that it always works. Modern Turkish has some twenty thousand words in common use, and a target with three consonants has hundreds of near neighbours. The question is never "is there a match" but "is there more of a match than a string that cannot mean anything would get?" We built seven lexicons and gave each its own null.

The modern side of the comparison uses open spelling dictionaries rather than a published dictionary's headword list, because a spelling dictionary can be counted. Turkish is the *hunspell-tr* word list, whose 292,000 forms are cut to 21,721 by keeping only words with a wordfreq Zipf score of at least 2.5, that is, words a speaker actually uses.^{25,23} Kazakh is the *myspell-kk* list, 53,182 forms, romanised on load by a fixed Cyrillic-to-Latin table.¹³ The frequency filter matters: without it the Turkish list is dominated by inflected and rare forms, and every target finds a match. The historical side is described in §3.4.

LEXICON	ENTRIES	PSEUDO-NAMES FINDING A MATCH	MEDIAN SCORE FOR A PSEUDO-NAME	NAMES BEATING CHANCE, $P \leq 0.05$
Turkish, common use	21,721	231 / 250	0.156	2 of 36
Kazakh	53,182	239 / 250	0.183	0 of 36
Union of Turkish and Kazakh	73,406	236 / 250	0.117	2 of 36
Turkish, loan-shaped words removed	14,876	233 / 250	0.200	2 of 36
Intersection (in both languages)	1,497	219 / 250	0.329	3 of 35
<i>Dīwān Lughāt al-Turk</i> , 1072–77	6,337	276 / 300	0.200	3 of 36
<i>Codex Cumanicus</i> , c. 1300	2,067	275 / 300	0.242	3 of 36

TABLE 3 A lower score is a closer match, so a *lower* median in column four means the lexicon is noisier: meaningless strings are finding better matches in it. Read that column as the price of admission. Two results are worth separating out. **Pooling two dictionaries makes matters worse:** the union finds more matches for the real names and more for the meaningless ones, and the noise floor falls from 0.156 to 0.117 while nothing else moves. A larger haystack does not help you find a needle. What helps is the opposite operation: keeping only words present in *both* languages lifts the floor to 0.329, because a word must survive two independent borrowing histories to qualify. That cure has its own disease: the intersection is over identical spellings, and cognates diverge, so Turkish *tanrı* and Kazakh *tāñır* both fall out while *domino* and *kauçuk*, borrowed into both recently, survive. It enriches for shared modern loans. Across all seven lexicons, between none and three names in thirty-six clear $p \leq 0.05$, which at a one-in-twenty threshold is what thirty-six coin flips give.

We can check two cases against a known answer, and the older lexicons do better on both.

NAME	READING	LEXICON	RANK	CANDIDATE	GLOSS IN THAT DICTIONARY
撐犁 chengli	ḍaŋ lei	Modern Turkish	—	not returned	<i>dengeli</i> "balanced" came first
		<i>Dīwān Lughāt al-Turk</i>	4th	tengri	<i>gök, sema</i> ("sky, heaven")
		<i>Codex Cumanicus</i>	2nd	tengri	<i>deus</i> ("God")
頭曼 Touman	do man ^c	Modern Turkish	—	duman	"smoke"; <i>tümen</i> not in the top three
		<i>Dīwān Lughāt al-Turk</i>	3rd	tümen	<i>pek çok</i> ("very many")
		<i>Codex Cumanicus</i>	2nd	touman	<i>nebula</i> ("mist"); <i>tümen</i> absent from this lexicon

TABLE 4 The *Hanshu* glosses 撐犁 as 天 "heaven". Kāšgarī in Baghdad and the compilers of the *Codex* in the Crimea, working a thousand years later with no knowledge of the Chinese record, gloss *tengri* as "sky, heaven" and as *deus*. The attestations are independent and the search connects them, which it fails to do against a modern dictionary, and does better the older and smaller the lexicon becomes. The second name repays attention because it cuts against the accepted etymology: *tümen* "ten thousand" is absent from the Cuman vocabulary entirely, and in every lexicon holding both, the search prefers a different lexeme: Turkish *duman*, Kāšgarī's *tuman*, Cuman *touman*, all "smoke, fog, mist", one word across three lexicons and a thousand years. The difference is a vowel, and §7 puts vowel recovery at 50%, so this settles nothing; it is recorded because the alternative is never raised.

A caution about the strongest datum. The match between *ᠳaŋ lei* and *tengri* is the most quoted evidence for a Turkic Xiongnu. The inference only works if *tengri* is Turkic to begin with. The word is not confined to Turkic: Mongolic has *tenger* "sky", Turkic *tengri* survives as modern Turkish *Tanrı*, and Georg has argued that the Turko-Mongolic word is built on Proto-Yeniseian **təŋgVr-* "high" with a Turkic suffix, making it a borrowing rather than an inheritance.^{8,28} That derivation is one argument among several and need not be accepted; what holds either way is that a word attested across Turkic, Mongolic and, on Georg's account, Yeniseian cannot discriminate between them. 撐犁 shows that the Xiongnu used, for "heaven", a word also found in Turkic and Mongolic. It does not show which of those languages they spoke.

7 Answering the 1974 question

The Xiongnu names survive in Han-period sources; the Turkic training corpus is Ming, twelve hundred years later. We can measure directly whether anything carries across that gap: we take characters whose readings are known in both layers, predict the earlier reading from the later one, and score by component.

PART OF THE SYLLABLE	RECOVERED CORRECTLY MING → LATER HAN	95% CONFIDENCE INTERVAL	CHANCE BASELINE: BEST BLIND GUESS	LIFT: POINTS GAINED OVER THAT GUESS	CLASS OF SOUND LAW ¹⁵
Nasal ending (-m, -n, -ng)	99.2%	95.4 – 99.9	35.4%	+63.8	deterministic
Initial consonant	79.3%	72.4 – 84.8	9.1%	+70.1	strong probabilistic
Vowel	50.0%	42.4 – 57.6	15.2%	+34.8	weak probabilistic
Stop ending (-p, -t, -k)	0.0%	0.0 – 7.9	35.4%	-35.4	no law

TABLE 5 $n = 164$ characters (nasal ending 118, stop ending 45). "Recovered correctly" is leave-one-out majority prediction of the earlier reading from the later; intervals are Wilson score intervals; the chance baseline is the unconditional majority class for that part, the score obtained by always guessing the commonest value. **Lift is column two minus column four:** how many percentage points knowing the later reading actually buys. Do not read an aggregate coda figure: blending nasals and stops gives a flattering 87% that means nothing, because the sample is 72% nasal-ending characters.

This is the reviewer's question, answered. Three of the four components classify cleanly. Nasal endings, at 117 of 118, are deterministic for practical purposes. Initial consonants, at 79.3% with the interval nowhere near one half, are a strong probabilistic law. Stop endings are not a law in any sense: not one of 45 survived, and the interval tops out at 7.9%.

The fourth exposes a flaw in the rule itself. Vowels land at **50.0%** (82 of 164, on the boundary to the decimal), and the confidence interval, 42.4 to 57.6, straddles it. By the reviewer's criterion the vowel cannot be classified, and no larger sample would help, because the estimate is not uncertain; it is genuinely at one half.

But one half is the right threshold only when there are two outcomes. A vowel is a choice among many, and the honest comparison is against what you would get knowing nothing, the majority class, here 15.2%. Against that, vowels carry **+34.8 points** of real information: a weak probabilistic law, not a coin flip. Stop endings, by the same measure, fall *below* their baseline: knowing the Ming reading makes you worse off than guessing, because Old Mandarin merged -p, -t and -k into nothing and the merger is not recoverable.

So the fifty-year-old question has a shape the fifty-year-old debate could not see. The threshold is not a constant. It is the baseline of the component in question, and it moves with how many outcomes that component has. That correction costs nothing to state and could not have been stated without counting.

8 Three measurements of scribal practice

8.1 Is a Chinese *m-* ever a foreign *b-*?

Reading 冒頓 as *Bayatur*, the most repeated etymology in the field, requires the *m-* to stand in for a *b-*. Proto-Turkic has no initial **m-*: the *m-* words in Old Turkic are loans, or come from **b-* assimilating to a nasal immediately after the vowel: **bän* → *men*, **buŋ* → *muŋ*.^{4,7} So on the Turkic reading the substitution is load-bearing. Each corpus is a set of transcriptions whose source word is known, which makes the question a counting exercise.

CORPUS	SOURCE WORDS BEGINNING B-	WRITTEN WITH A CHINESE M- CHARACTER
Mongolian (<i>Secret History</i>)	738	0
Turkic (Ligeti glossaries)	78	0
Sanskrit (Later Han, verified)	22	3
Total	838	3 (0.4%)

TABLE 6 The reverse direction is clean: Chinese m- characters render a source m- in 303 of 305 Mongolian cases, 90 of 97 Sanskrit and 14 of 17 Turkic. The forward direction, which is the one that matters here, must be read by layer and not pooled. The total of 3 in 838 is 0.4%, but all three instances fall in the *Later Han* corpus, where the rate is **3 of 22, or 13.6%**, with a Wilson interval of 4.7 to 33.3. The two later corpora contribute 0 of 816 between them. Pooling across eight centuries of Chinese produces a figure that describes no period at all, which is the point §7 makes in general terms.

Read by layer, the substitution is not exceptional. A rate of 13.6% is the same order as ts- for a source affricate, which §6.3 treats as a minority but attested spelling at 15%. The scribe did have the alternative: within the Later Han corpus a source b- took a b- character 26 times in 32. But having an alternative and always using it are different things, and one time in seven he did not. **So an m- character is not, by itself, an argument against a Turkic reading of 冒頓.**

The onomastic record points the same way, in a later period. Tang scribes wrote the Türk ruler known in the inscriptions as Qapaġhan Qaġhan with the personal name 默噶, read as *Bögü-Çor*; the Uyghur ruler 牟羽 is the same word *bögü* "wise" behind a different m- character; and 木桿 stands for *Muqan*. Two different Chinese m- characters for one Turkic b- word is a pattern, not an accident. These are Tang rather than Han, and Late Middle Chinese treated its nasal initials differently, so they corroborate that the practice existed rather than fixing its Han-period rate; but they remove any sense that the substitution would be unheard of in a ruler's name.

The objection moves to the other end of the word. *Bayatur* ends in -r, and 頓 *tuən*^c carries a written -n. Later Han has no -r coda, so a scribe facing a source syllable such as *tur* had to choose something, and the Sanskrit corpus records what he chose. Of 37 source syllables ending in -r, he left the syllable **open** 21 times (57%), used a -t coda **14** times (38%): 達 for -har in *bodhidharma*, 羯 for kar- in *ākarṣaṇī*. A -p coda occurs once, and a -n coda **once**, 3%. Reading 頓 as *tur* therefore asks for the rarest of the four options, where the expected spelling was an open character or a -t.²² The syllable count is no obstacle to either reading. *Bayatur* has three syllables against two characters, but 93 of the 1,017 verified pairs use fewer characters than syllables, and in 47 of those a character with a written stop coda does the compressing. 竭 covers -gata in 多阿竭 for *tathāgata*, 塞 covers -saka in 優婆塞 for *upāsaka*. Compression is a real device, available to any reading of 冒頓, so the final -r stands alone as the difficulty.

What survives on the Turkic side is unaffected by any of this. Proto-Turkic has no initial *m-, and the one route to one is *b- assimilating to a nasal immediately after the vowel; 冒 *mək* has no nasal in that position, its only nasal being the -n of 頓. So on a Turkic reading the m- has to be the scribe's choice rather than the source language's sound, and the count above says he was entitled to make it. Small numbers should be kept in view: 3 of 22 is four cases short of nothing, and the interval is wide.

8.2 How was a source coda spelled?

A related question, and the one §9 turns on. When a source syllable ended in a consonant, did the scribe use a character ending in that consonant, or an open character with a vowel after it? The Sanskrit corpus answers it for the right period: for the 770 pairs aligning one character to one syllable, each Sanskrit syllable is matched to its character and the coda of each recorded.

SOURCE SYLLABLE ENDS	N	CLOSED WITH THE SAME CODA	WRITTEN WITH AN OPEN CHARACTER
-n	93	83%	13%
-ŋ	22	73%	9%
-t	53	53%	28%
-k	21	48%	38%
-p	10	20%	20%
All stops	84	48%	30%

TABLE 7 Roughly a third of stop-final source syllables were written with an open character, against 13% and 9% for nasals. An open spelling after a stop is ordinary, not exceptional. Note the shape of the asymmetry: nasals are handled robustly and stops loosely, the same split Table 5 found surviving the period gap, appearing here in what the scribe chose to write rather than in what survived. Two independent measurements of one underlying fact.

8.3 When a consonant is written twice, was it said twice?

A scribe could spell one source consonant across a syllable boundary: the first character ends in *-t*, the second begins with *t-*. Whether that doubling records a geminate or is a convention of the script changes how many consonants a transcription actually claims, and therefore how a name may be read. It is answerable on the Sanskrit corpus, where 772 pairs align one character to one syllable.

We take every junction at which a written stop or nasal coda is followed by a same-place onset, and ask whether the Sanskrit form carries that consonant once or twice:

JUNCTIONS WHERE A WRITTEN CODA IS FOLLOWED BY AN ONSET OF THE SAME PLACE	THE SOURCE WORD HAS THAT CONSONANT ONLY ONCE	THE SOURCE WORD HAS IT TWICE
140	121 (86%)	19 (14%)

TABLE 8 The double spelling is a convention, not a geminate: 釋迦 for *śākya*, 怛他 for *tathā*, 富單 for *pūtana*. A sequence such as 骨都 *kuət ta* therefore transcribes a source *kuta* with one *t*, not *kutta*. The count is conservative: "twice" is scored whenever the doubled letter appears anywhere in the Sanskrit form, so a geminate elsewhere in the word counts against the merger, and 86% is a floor rather than a ceiling. Like §8.2, this is a measurement that settles an objection which would otherwise be argued from intuition: a reader who assumes two written consonants mean two spoken ones will reject readings that the corpus permits. The nasal case is cleaner still: every source word in the corpus with a medial *ŋg* or *ŋk* is written with a *-ŋ* coda followed by a velar onset, 13 of 13: 央具梨 for *aṅguli*, 商羯羅 for *śaṅkara*, 僧伽施 for *sāṅkāśya*. Two characters spelling one medial nasal is not a licence the scribe occasionally took; it is the only device available to him, since a Chinese syllable cannot begin with *ŋ*.

9 A worked example: what a proposal must survive

Section 6.3 shows that reading a transcription against a dictionary is, on its own, at chance. That is not an argument that no such reading can ever be supported; it is an argument that support must come from somewhere other than the resemblance. One case in this material meets that standard, and it is set out here as a worked example rather than as a result about the Xiongnu.

From 呼韓邪 onward, every chanyu title ends in 若鞮, Later Han *ńak te*. The *Hou Hanshu* states that the element was added to the title from that reign and glosses it 孝 "filial". This is the only Xiongnu title element the Chinese sources translate, which means the meaning is *recorded*, not inferred from the sound.

Four things then have to hold, and each is checkable independently of the others.

- **The sense must be attested.** We searched the two lexicons of §3.4 and found *Inak*, glossed *fidelis* "faithful", in the *Codex Cumanicus*.¹⁴ It is absent from Kāšgarî, so the word belongs to the Kipchak tradition. Old Turkic *inaq* "trusted one, confidant" is generally derived, with *inan-* "to trust", from a root *ina-* "to rely on". The sense the proposal needs is documented in a medieval glossary, and it matches 孝 independently.
- **The word is of a kind that becomes a title.** This matters more than the gloss alone. Clauson records Chagatay *inak* as *nāyib va muqarrab*, "a royal representative or senior minister", and the cognate *inanç*, from the same *ina-* root, appears in Old Turkic as an office and as a name element: *él inanç tirek* in a Uyghur list of proper names, *inançları buyrukları* "his ministers and officers", *inanç tayanç* for a confidential minister to a *beg*, and *él ögesi inancu bilge* in the Yenisei epitaphs of southern Siberia.⁴ A court title built on this root is therefore not a stretch imposed on the word; it is what the word

ordinarily does. The *Hou Hanshu* has 若鞮 entering the formal name from a particular reign onward, which is exactly how such a title spreads.

- **The morphology must be licensed.** A reading *inaq-ti* is not available: *-ti* is not a noun-forming suffix, and the lexicons show it rather than a grammar having to assert it. Of the 31 Cuman headwords ending in *-ti*, almost all are glossed with Latin perfects (*Ayti: dixit, Azti: transgressus est, Öwtiti: deterisit*), the definite past ending; the remainder are numerals or Arabic loans. Kāşgarî's index has six, none containing a suffix. A title cannot be a past-tense verb. Old Turkic does have a suffix that fits: *-t*, a plural used specifically with titles: *tegin* → *tegit*, *tarqan* → *tarqat*.⁷ On that reading the element is **Inakt**, "the faithful ones", and 若鞮 is not one ruler's epithet but a title borne by a line, which is exactly how the Chinese sources use it.
- **Objections to the spelling must be tested, not asserted.** Against a final *-t* it can be objected that Later Han had *-t* final syllables and a scribe would have used one rather than the open 鞮 *te*. That is a claim about scribal habit, and Table 7 counts it: 30% of stop-final source syllables were written with an open character. The objection does not hold.
- **What remains unexplained must be stated.** If the source was *inakt*, the opening *t-* is not written, and 若鞮 has no room for it. That objection stands.

Three of four hold, with the fourth named. That is the standard this paper asks of a proposal, and it is deliberately higher than resemblance: the sense is attested in a source contemporary with neither party to the comparison, the morphology is licensed by a suffix with the right distribution, and the one objection that could be tested was tested and failed.

One further reading of our own, 冒頓 as *Bögü Tın*, does meet this standard, and §9.2 sets it out. We produced 31 others in the course of this work. None of them meets it (by §6.3 they are at the chance rate) and they are therefore **not presented here as findings**. We include them as **Supplementary Material S1** and release them in the repository as [author_proposals.csv](#), dated and labelled, because readers will run this search in their heads regardless and it is better that the output is written down than imagined.²⁶

We supply S1 for the record, not for review. It has the status of a laboratory notebook that happens to be legible: each row gives the Chinese form, the Later Han reading, the author's rendering, the sense proposed for it, and, where one exists, the historical or onomastic parallel that suggested it. Those parallels are drawn mostly from later Turkic naming practice, where a ruler's epithet is a common noun (Bayezid *Yıldırım* "the thunderbolt", Alp Arslan "the lion", Mehmed *Avcı* "the hunter"), and they are offered as the kind of thing that would have to be argued rather than as an argument. A reader who accepts none of them loses nothing from the paper; the measurements in §§5–8 do not rest on any of them. The one reading that is argued here, *Inakt*, is in S1 too, marked as the exception, so that the file can be read as a single list with the standard visible in it. S1 also contains a worked demonstration, at §S1.5, of why the decoder's own output must not be quoted in support of an individual reading, the concrete form of the warning in §6.2.

9.1 Two cases which are not ours to propose

We can apply the same standard to comparisons somebody else has already made, and that is the more useful exercise, because the reading is then not ours to want: the selection effect §6.3 measures is removed by construction. Two items in the record carry a reading from the literature and a Chinese gloss to check it against.

The first is 徑路 *keŋ^c la^c*, the 徑路刀, a Xiongnu blade used in oath-taking. Friedrich Hirth compared it with Turkic *kıŋrak* in 1908, calling the word the oldest Turkish on record.¹² Kāşgarî glosses *kıŋrak* as a large knife resembling a cleaver, for cutting meat and dough. The Chinese side independently says the referent is a blade, two sources with no contact agreeing about what the object is, which is a stronger kind of agreement than a sound match. The consonant frame is exact: 徑 gives the velar onset with its *-ŋ* coda written, and 路 gives a Chinese *l-* for a source *r-*, the direction attested in 12 of 14 cases. The Turkic word's final *-k* is not written, and Table 7 gives the rate at which a stop-final source syllable was spelled with an open character: 30%.

Where the Turkic word comes from. This is not left to inference. Clauson derives *kıŋra:k* as a deverbal noun from an unattested **kıŋra:-*, itself denominal from *kıŋır*, in the sense of something curved or something that cuts crookedly.⁴ Kāşgarî's own index supports that family on the same frame: *kıŋır*, *kıŋru* "askew, squint", *kaŋrak* "palate", the arch of the mouth, and the verb *koŋur-* "to wrench, bend back", surviving as Turkish *kanırmak*; and the *Codex Cumanicus* independently glosses *Kingir* as *curvus*. A curved blade named from bending is therefore motivated inside Turkic by the standard authority, not by us.

The word also survives. Clauson records it in Teleut as *kıırak*, and in Kirghiz as *kııarak* / *kııırak*, "a rough two-edged knife used for cutting felt, scraping hides and sheepskins", the same object, in living languages, a thousand years after Kāšgarî and two thousand after the *Hanshu*.

Two cautions belong with this. Kāšgarî's text at III 382 reads *kırpa:k*, corrected in the margin to *kııra:k*, so the form rests on a marginal emendation rather than the text as written. And of the family it is *kıırak* that fits the transcription exactly, at three consonants with nothing supplied; *kıırak*, the member that means a knife, needs the unwritten final *-k* that Table 7 licenses at 30%. The form that fits best does not mean "blade", and the form that means "blade" does not fit best.

And the rival account. The same object has been derived instead from an Iranian word for a short sword, the type whose Greek name is *akinakes*. The history supports it: the Xiongnu's western neighbours were Iranian-speaking, and the weapon is present in steppe archaeology, so a borrowed name is not a stretch.

Unusually, that account is testable with the same machinery, and it does not survive as well. 路 *la*^c is a liquid-onset character, so the channel requires a liquid in the second position of the source word. The *akinakes* family (*k-n-k*) has no liquid anywhere in it. If the proposal is that 徑路 transcribes that word, it fails on the frame, not at the margins.

Two cautions against reading this as a verdict. We have not seen the exact Iranian etymon its proponents advance, and a different Iranian word could fit where *akinakes* does not; the test applies only to the form named. And our channel establishes what a transcription permits, never what it requires: a Turkic reading surviving does not make a Turkic source the only one available. This is Doerfer's objection in its sharpest form, and the honest report is that one candidate survives its test and the other, as stated, does not.

The second is 屠耆 *da gi*, which the Chinese sources gloss 賢 "worthy, wise" and which appears in 左右屠耆王, the Wise Kings of the Left and Right. The Turkic-hypothesis literature reads it **toyri*. The *Codex Cumanicus* has *Togru* glossed *rectus*, and the skeleton matches at the same single step, a Chinese *l*- for a source *r*-.

The semantic fit is closer than "straight" alone suggests. Clauson derives *toğuru*: / *toğru*: as a gerund of *toğur*-, physically "straight" and then metaphorically "straight, honest, upright, true", and that moral sense is old rather than modern: a fourteenth-century glossary renders it with Arabic *al-tiqa* "trustworthy, honest", and *toğru ayt*- with *şadaqa* "to tell the truth".⁴ The word survives across the whole family and as modern Turkish *doğru*. So the comparison is not between 賢 "worthy" and a word for physical straightness, but between 賢 and a word whose ordinary sense is "upright, honest, trustworthy".

The root family also yields names, which is the argument §9 makes for *inak*: Kāšgarî records that men are called *Toğrı* after a bird of prey, an *Alp Toğrul Tégin* appears in a Uyghur list of proper names, and the first Seljuk sultan bore the name. What remains soft is the vowel, and the fact that 144 attested headwords fit this frame, so the transcription does little to narrow the choice.

So two existing comparisons survive tests that 32 of our own readings fail, and one of the two has a rival etymology our method cannot rule out. That asymmetry is the point of this section. The channel cannot generate supported readings at this data density, but it can adjudicate ones already on the table, and adjudicating other people's proposals is a use of it that the accuracy figures actually license.

9.2 A reading the measurements do support: 冒頓

The counts of §8.1 constrain a reading of 冒頓 from both ends, and taken together they point somewhere. 冒頓 *mək tuən*^c is the most consequential name in the record, and the frame it fixes is unusually tight, because **both codas are written**: the *-k* of 冒 and the *-n* of 頓 are on the page and have to correspond to something. The frame is therefore *b/m*, *k/g*, *t/d*, *n*: four consonants, ending in a nasal.

The first character has a documented Turkic value. This is not an inference from a rate. In the Later Han table 冒 and 默 have the *same* reading, *mək*. And 默噉, the personal name of the Türk ruler known in the inscriptions as Qapağan Qaghan, is read *Bögü-Çor*. So that exact Later Han syllable is on record writing *bögü* in a Turkic ruler's name: the *m*- standing for the source *b*-, the written *-k* standing for the source *g*, and the final vowel of *bögü* not written at all.

Each step of that is separately licensed by the corpora. The *m*- for a source *b*- is 3 of 22 in the Later Han layer, 13.6% (§8.1). One character covering two source syllables, with a written stop coda doing the compressing, occurs in 47 of the 1,017 verified pairs: 竭 for *-gata* in *tathāgata*, 薩 for *-sattva* in *mahāsattva*, 弗 for *-putra* in *śāriputra*. 冒 for *bögü* is that same shape.

The second character has attested candidates and needs no licence at all. 頓 *tuən*^c is a dental onset with a written *-n*, and the two lexicons supply *tın* "spirit, breath" (Codex *Tin*, *anima*), *tün* "night" (Codex *Tun*, *nox*), *tan* "dawn" and *tun* "rest". Every one of them matches the frame with nothing inserted and nothing discarded.

So we read 冒頓 as *Bögü Tın*, "sage spirit", and prefer *tın* among the four because of what *bögü* is. Kāšgarî glosses it *bilgin*, *akıllı*, *hakim*; Clauson says the word "seems to connote both wisdom and mysterious spiritual power", and it was borrowed early into Mongolian as *böge*, the ordinary word for a shaman. It is used in regnal styles across the record: *Teŋri Bögü Teŋriken*, Bögü Qaghan of the Uyghurs, Bögü-Çor of the Türks. A name pairing that word with the word for spirit is coherent for the founder of a steppe confederation in a way that no dictionary lookup produced; it fell out of the frame.

Measured against the standard of §9. The sense of both elements is attested independently of any Chinese source. The morphology is the commonest Turkic name shape, two nouns in apposition, as in Alp Arslan or Bögü-Çor itself. The objections were tested rather than asserted, and the one that would have killed it, a Chinese *m-* for a foreign *b-*, was measured and dissolved. And what remains unexplained is stated: **the compound *bögü tın* is not itself attested**, only its two halves; the vowel of the second element is not recoverable, so *tün*, *tan* and *tun* are equally available and choosing *tın* is a semantic preference rather than a phonological result; and the tone category of 頓 is not used.

Set beside it, *Bayatur*, the field's reading for a century, fails on the one point that the corpus decides cleanly: it ends in *-r* against a written *-n*, a spelling that occurs once in 37 opportunities, where the scribe's normal choices were an open character or a *-t*. Our reading asks for no correspondence rarer than one in seven, and asks nothing at all of the final consonant.

We are aware of what this section is. With §9 it is one of only two places in this paper where a Xiongnu name is read positively, and the temptation to overstate such a reading is exactly the one §6.3 measures. So the claim is bounded: *Bögü Tın* is a reading that survives every test we can apply, on the tightest frame in the record, and it is offered against a standard etymology that does not. Whether it is right is a question for Turcology, and the material for deciding it, the counts and the corpora and the scripts, is in the repository.

10 Where the method does work

The constraint identified in §5 is data density, not the difficulty of the language. That suggests where to go next. Göktürk-era material has what the Xiongnu material lacks: a source language that is not in dispute, an independent script recording it, and enough of it to train on.

CHINESE, AS WRITTEN IN THE TANG HISTORIES	LATER HAN READING OF THOSE CHARACTERS (SCHUESSLER)	THE SAME READING IN TURKISH ORTHOGRAPHY	INDEPENDENTLY IDENTIFIED AS, IN ESTABLISHED SCHOLARSHIP
葉護	jap ɣua ^c	yap ğua	<i>yabghu</i> , a princely title
可汗	kʰai ^b gan ^c	kai gan	<i>qayan</i>
特勤	dək giən	dık gıın	<i>tigin</i> , prince
突厥	tʰuət kyat	tuıt kyat	<i>Türk</i>
阿史那	ʔa ɕiə ^b na	ʼa ɕiı na	<i>Ashina</i> , the ruling clan
黠戛斯	geɬ keɬ sie	get ket sie	<i>Qırqız</i> (Kyrgyz)

TABLE 9 Neither middle column is model output. Column two is Schuessler's published Later Han reconstruction applied to Tang-era words, deliberately the wrong period, which is the point of the exercise. The last column is not an output of this work either: each identification is long settled on evidence independent of the Chinese transcriptions, the Orkhon inscriptions giving *qayan*, *tigin*, *yabghu* and *Türk* directly in Old Turkic script, several corroborated again in Sogdian, Arabic and Byzantine sources.^{7,9} They function here as ground truth. That 葉護 → *yap ğua* is readable against *yabghu* despite a period mismatch of four centuries measures the conservatism of Chinese scribal practice, not linguistic continuity between the Xiongnu and the Göktürks, on which this material bears not at all.

Chinese records name a great many steppe peoples, and they fall into three groups for these purposes. **Tractable**: those recorded in the Tang histories whose language is known: 鐵勒 Tiele, 回紇 Uyghur, 葛邏祿 Qarluq, 突騎施 Türgesh, 黠戛斯 Kyrgyz, alongside the Göktürk material. These have transcriptions, a securely identified source language and published readings for the right layer; they are training data. **Readable but uninformative**: groups such as 丁零 Dingling and the Avars, whose Chinese readings can be recovered but whose language is unknown, so a reading cannot be checked. **Out of reach**: the Xiongnu themselves, and the

Pechenegs, who are barely in the Chinese record at all. The distinction is not how interesting the group is but whether there is anything to validate against.

11 Limitations

- **No independent philological review.** The computational method and its evaluation are ours to defend. The philology is not our work, and we have not had it checked for this paper: the Later Han readings are Schuessler's, the etymologies discussed are from the published literature, and the identifications in Table 9 are the established consensus. A specialist reading would strengthen the paper and might revise judgements in §6 and §9.
- **Corpus validation is internal.** The extraction rules are documented and internally consistent (93.4% length agreement for Mongolian, an induced akṣara inventory matching the standard one for Sanskrit), but no external reviewer has checked a sample against the sources.
- **The Turkic corpus is Ming-era,** roughly 1,200 years later than the Xiongnu material. §7 measures what that costs rather than assuming it away.
- **Reconstructions are inferences.** Schuessler, Pulleyblank and Ning are three scholars with three notations; some apparent disagreement between them is notational rather than substantive.
- **Names travel.** Even a perfect reconstruction would not settle ethnicity. Turkic tribes bore Iranian names and Mongol rulers used Turkic titles; the caution in §6.3 about *tengri* is a specific case of this general problem.
- **Reproducibility has a version dependency.** The headword count for the Kāṣṣgarī lexicon differs between machines, 6,820 against 6,880, according to the version of the PDF text-extraction library. It shifts the noise floor by 0.008 and changes no rank or conclusion, but it is the kind of thing that stays invisible until someone else runs the code.

12 Data and code availability

We package everything described here as one repository, so that anyone can regenerate every number and table in this paper from the raw inputs by running the numbered scripts in order. The repository is the durable output; this paper describes it.²⁶

```
chinese-transcription-channel/
├─ README.md           what each output is, and how to regenerate it
├─ LICENSE              code MIT
├─ DATA_LICENSE.md    the three-way split described below
├─ data/derived/        the corpora, the lexicons, and every measurement
├─ src/                 steps 1–23 plus lookup_lexicons.py, each rerunnable alone
├─ reports/             one summary per step, with the numbers quoted above
└─ docs/                design notes, including approaches that were rejected
```

Supplementary Material S1. *Candidate readings of the Xiongnu record* accompanies this submission and carries the same content as [data/derived/author_proposals.csv](#) in the repository. We supply it for the record under §9 and do not offer it as a result of this paper.

Released, CC BY 4.0. The Mongolian and Sanskrit corpora, the Cuman lexicon (Kuun's 1880 edition being public domain), the induced akṣara table, the Xiongnu record of Appendix A, and every measurement file behind Tables 1–9.

Withheld, with the extraction script released instead. Four derived files inherit restrictions from their sources: the Turkic pairs, from Ligeti's glossaries (© Akadémiai Kiadó); the Kāṣṣgarī lexicon, from a modern Türk Dil Kurumu edition; and two files built on that lexicon: the Old Turkic echo results and the candidate-space listing, the latter also quoting Clauson's *Etymological Dictionary* verbatim in a context column. Anyone holding the sources can regenerate any of them byte-for-byte by running the numbered script, which keeps the paper's numbers checkable without redistributing the source. No permission is required for this arrangement and none has been sought.

Not in the repository. Third-party source texts and datasets (dictionaries, scanned volumes, Unihan, the DILA authority files) are inputs, not outputs. Where a script needs one, its docstring says which file and where to obtain it.

A Appendix: the Xiongnu record

All thirty-nine names, titles and lexical items in the Chinese histories, with the Later Han reading of each character and a Turkish-orthography rendering for readers who do not read phonetic notation. The Later Han column is published scholarship, not model output; our contribution is the assembly: we collected the names in one place and joined the readings to them character by character. Superscript letters mark Han tone categories. 若鞮 is listed once, shaded, rather than repeated in the six titles that carry it.

A.1 Rulers and the ruling clan

NAME OR CLAN NAME, AS WRITTEN IN THE HISTORIES	LATER HAN READING OF THOSE CHARACTERS (SCHUESSLER)	THE SAME READING IN TURKISH ORTHOGRAPHY	DATES OF RULE, AND THE SOURCE IN WHICH THE FORM APPEARS
頭曼 Touman	do man ^c	do man	209 BCE
冒頓 Modu / Modun	mæk tuən ^c	mık tuın	209–174 BCE
稽粥 Jizhou / Jiyu	k ^h ei ^B tśuk	kei çuk	174–160 BCE
軍臣 Junchen	kun dźin	kun cin	160–126 BCE
伊稚斜 Yizhixie	ʔi ɖi ^c ja	’i di ya	126–114 BCE
烏維 Wuwei	ʔa wi	’a vi	114–105 BCE
詹師廬 Zhanshilu	tśam ʃi lia	çam ʃi lia	105–102 BCE
响犁湖 Xulihu / Goulihu	huo lei ga	huo lei ga	102–101 BCE
且鞮侯 Qiedihou	tsa ^B te go	tsa te go	101–96 BCE

NAME OR CLAN NAME, AS WRITTEN IN THE HISTORIES	LATER HAN READING OF THOSE CHARACTERS (SCHUESSLER)	THE SAME READING IN TURKISH ORTHOGRAPHY	DATES OF RULE, AND THE SOURCE IN WHICH THE FORM APPEARS
狐鹿姑 Hulugu	ɣua lok ka	ğua lok ka	96–85 BCE
壺衍鞬 Huyandi	ga jan ^B te	ga yan te	85–68 BCE
虛間權渠 Xulüquanqu	k ^h ia lia gyan gia	kia lia gyan gia	68–60 BCE
握衍胸鞬 Woyanqudi	ʔɔk jan ^B guo te	’ok yan guo te	60–58 BCE
呼韓邪 Huhanye	ha gan ja	ha gan ya	58–31 BCE
郅支 Zhizhi	t́sit t́se	çit çe	56–36 BCE
若鞬 Ruodi	ńak te	nyak te	title element
復株累若鞬 Fuzhulei Ruodi	buk ʈio lui ńak te	buk tio lui nyak te	31–20 BCE
搜諧若鞬 Souxie Ruodi	so ɣei ńak te	so gei nyak te	20–12 BCE
車牙若鞬 Cheya Ruodi	t́sʰa ɲa ńak te	ça nga nyak te	12–8 BCE
烏珠留若鞬 Wuzhuliu Ruodi	ʔa t́so liu ńak te	’a ço liu nyak te	8 BCE–13 CE
烏累若鞬 Wulei Ruodi	ʔa lui ńak te	’a lui nyak te	13 CE–18 CE
呼都而尸道皋若鞬 Huduershi Daogao Ruodi	ha ta ńə sí dou ^B kou ńak te	ha ta nyı şı dou kou nyak te	18 CE–46 CE
攣鞬 Luandi	lyan te	lyan te	–
虛連題 Xulianti	k ^h ia lian ^B de	kia lian de	–

Twenty-one *chanyu* from Touman to Punu, plus both forms of the ruling clan name: 攣鞬 in the *Shiji* and *Hanshu*, 虛連題 in the *Hou Hanshu*.

A.2 Titles and lexical items

TERM, AS WRITTEN IN THE HISTORIES	LATER HAN READING OF THOSE CHARACTERS (SCHUESSLER)	THE SAME READING IN TURKISH ORTHOGRAPHY	WHAT KIND OF ITEM IT IS: TITLE, LEXICAL WORD, OR ETHNONYM
單于 chanyu	d́zan wa	can va	title
撐犁孤塗單于 chengli gutu chanyu	ɖaŋ lei kua da d́zan wa	dang lei kua da can va	title

TERM, AS WRITTEN IN THE HISTORIES	LATER HAN READING OF THOSE CHARACTERS (SCHUESSLER)	THE SAME READING IN TURKISH ORTHOGRAPHY	WHAT KIND OF ITEM IT IS: TITLE, LEXICAL WORD, OR ETHNONYM
撐犁 chengli	ḍaŋ lei	dang lei	lexical
孤塗 gutu	kua da	kua da	lexical
閼氏 yanzhi / ezhi	ʔat dʒe ^B	ʼat ce	title
屠耆 tuqi	da ɣi	da ɣɪ	title
谷蠡 luli / guli	kok le ^B	kok le	title
當戶 danghu	taŋ ga ^B	tang ga	title
骨都侯 guduhou	kuət ta go	kuɪt ta go	title
且渠 juqu	tʂa ^B ɣia	tʂa ɣɪa	title
徑路 jinglu	keŋ ^c la ^c	keng la	lexical
服匿 funi	buk ɳik	buk nɪk	lexical
甌脫 outuo	ʔo tʰuat	ʼo tuat	lexical
胡祿 hulu	ɣo lok	go lok	lexical
匈奴 Xiongnu	huoŋ na	huong na	ethnonym
羯 Jie	kɪat	kɪat	ethnonym

The lexical items matter most: they are the only Xiongnu words the Chinese sources gloss. 撐犁 is glossed "heaven" and 孤塗 "son", which is why that pair has carried so much of the argument. Note that the Yeniseian proposal for 孤塗 is *(h)ala* "son" and the Turkic one *qut* "fortune"; the Chinese gloss 子 "son" fits the former and not the latter.^{27,28}

B References

1. Baley, D. (2025). *Chinese Transcription of Buddhist Terms in the Late Hàn Dynasty*, v7. Zenodo. doi:10.5281/zenodo.7905399. Data paper: *Journal of Open Humanities Data*, doi:10.5334/johd.110

2. Baxter, W. H. & Sagart, L. (2014). *Old Chinese: A New Reconstruction*. Oxford University Press.

3. Bonmann, S. et al. (2025). Linguistic evidence suggests that Xiōngnú and Huns spoke the same Paleo-Siberian language. *Transactions of the Philological Society*. doi:10.1111/1467-968X.12321

4. Clauson, G. (1972). *An Etymological Dictionary of Pre-Thirteenth-Century Turkish*. Oxford: Clarendon Press.
5. Doerfer, G. (1973). *Lautgesetz und Zufall: Betrachtungen zum Omnicomparatismus*. Innsbrucker Beiträge zur Sprachwissenschaft 10. Innsbruck.
6. Doerfer, G. (1973). Zur Sprache der Hunnen. *Central Asiatic Journal* 17: 1–50.
7. Erdal, M. (2004). *A Grammar of Old Turkic*. Leiden: Brill.
8. Georg, S. (2001). Türkisch/Mongolisch *tengri* 'Himmel, Gott' und seine Herkunft. *Studia Etymologica Cracoviensia* 6: 83–100.
9. Golden, P. B. (1992). *An Introduction to the History of the Turkic Peoples*. Wiesbaden: Harrassowitz.
10. Hú Zhènghuá & Huáng Rùnhuá (1984). 《高昌館雜字：明代漢文回鹘文分類詞匯》. Běijīng: Mínnú chūbǎnshè.
11. Hirth, F. (1908). *The Ancient History of China to the End of the Chóu Dynasty*. New York: Columbia University Press. (Comparison of 徑路 with Turkic *kırarak*, pp. 65–70.)
12. Kâşgarlı Mahmud (1072–77). *Dîwân Lughât al-Turk*. Index volume (*Dizin*), Türk Dil Kurumu Yayınları.
13. Kergalym (2021). *myspell-kk*: a Myspell/Hunspell spelling dictionary for Kazakh [dataset]. github.com/kergalym/myspell-kk. 53,182 headwords; romanised on load.
14. Kuun, G., ed. (1880). *Codex Cumanicus Bibliothecæ ad Templum Divi Marci Venetiarum*. Budapest: Editio Scient. Academiæ Hung. Public domain.
15. [Unsigned review] (1974). Lautgesetz und Zufall [review of Doerfer 1973]. *Anthropos* 69: 263–271. JSTOR 40458521.
16. Lee, J. et al. (2023). Genetic population structure of the Xiongnu Empire at imperial and local scales. *Science Advances* 9(15). doi:10.1126/sciadv.adf3904
17. Ligeti, L. (1966). Un vocabulaire sino-ouïgour des Ming. *Acta Orientalia Academiae Scientiarum Hungaricae* 19: 117–199, 257–316.
18. Ligeti, L. (1969). Glossaire supplémentaire au vocabulaire sino-ouïgour du Bureau des Traducteurs. *Acta Orientalia ASH* 22: 1–49, 191–243.
19. Pulleyblank, E. G. (1962). The consonantal system of Old Chinese. *Asia Major* 9: 58–144, 206–265.
20. Pulleyblank, E. G. (1991). *Lexicon of Reconstructed Pronunciation in Early Middle Chinese, Late Middle Chinese, and Early Mandarin*. Vancouver: UBC Press.
21. Savelyev, A. & Jeong, C. (2020). Early nomads of the Eastern Steppe and their tentative connections in the West. *Evolutionary Human Sciences* 2, e20.
22. Schuessler, A. (2009). *Minimal Old Chinese and Later Han Chinese*. Honolulu: University of Hawai'i Press. Data release: Zenodo, doi:10.5281/zenodo.14567004
23. Speer, R. (2022). *wordfreq*: word frequency data for 40+ languages, v3 [software]. Zenodo, doi:10.5281/zenodo.7199437
24. Tan, Y. Y. (2019). *Consonant Endings of Ming Dynasty Mandarin as Reflected in the Chinese Transcriptions of Uyghur Vocabulary*. MA thesis, University of Washington.
25. tdd-ai (2020). *hunspell-tr*: a Hunspell spelling dictionary for Turkish [dataset]. github.com/tdd-ai/hunspell-tr. 292,000 forms, reduced to 21,721 by the frequency filter.
26. Uluğ, A. M. (2026). *chinese-transcription-channel*: corpora, lexicons, channel models and evaluation code [software and data]. GitHub: github.com/ulug2004/chinese-transcription-channel
27. Vovin, A. (2000). Did the Xiong-nu speak a Yeniseian language? *Central Asiatic Journal* 44(1): 87–104.
28. Vovin, A. (2003). Did the Xiongnu speak a Yeniseian language? Part 2: Vocabulary. In *Altaica Budapestinensia MMII*, Proceedings of the 45th PIAC, Budapest, 389–394.
29. Wittek, A. et al. (2020). ASJP Database v19. Zenodo, doi:10.5281/zenodo.3843469